

LIS Nepal

Lalitpur, Nepal



EFK Stack: Documentation

Documented by

Aayam Adhikari
Software Intern - 2024

1. What is the EFK stack?

EFK, which stands for Elasticsearch, Fluentd, and Kibana, is a popular open-source stack used for collecting, processing, and visualizing logs.

Elasticsearch is a distributed search and analytics engine that can store and index large amounts of data. Fluentd is a data collector that can gather data from a variety of sources and send it to Elasticsearch for indexing. Kibana is a data visualization tool that allows users to interact with data stored in Elasticsearch and create custom dashboards and visualizations. Together, these three components provide a comprehensive logging solution that can handle a large amount of data and provide real-time insights into the behavior of complex systems. Below is a brief intro to each of these:

Elasticsearch: Elasticsearch is a distributed search and analytics engine designed for storing and indexing large volumes of data. It provides fast search capabilities and is highly scalable, making it ideal for use cases such as logging, search, and analytics within the EFK stack.

Fluentd: Fluentd is a versatile data collector that gathers logs from various sources and sends them to Elasticsearch for indexing. It supports a wide range of inputs and outputs, processes logs in real-time, and can enrich logs with metadata before sending them to Elasticsearch.

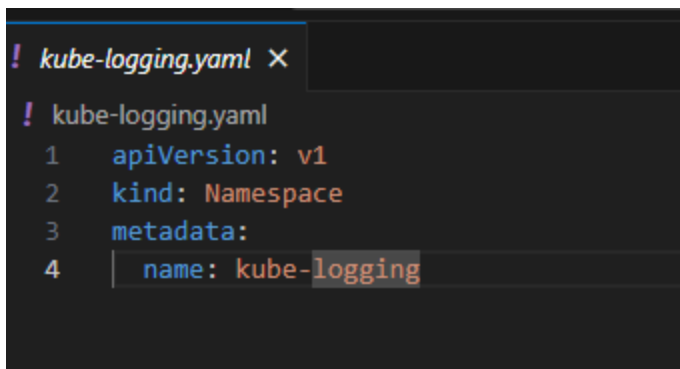
Kibana: Kibana is a powerful data visualization tool that allows users to interact with data stored in Elasticsearch. It enables users to create custom dashboards, charts, and visualizations to analyze and monitor logs in real-time. Kibana is essential for gaining insights into system behavior and monitoring application performance within the EFK stack.

2. Steps to set-up EFK stack for your logging solution

Step-1: Creating a Namespace

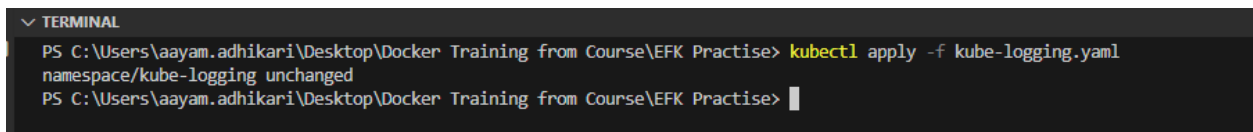
Before we roll out an Elasticsearch cluster, we'll first create a Namespace into which we'll install all of our logging instrumentation. Kubernetes lets you separate objects running in your cluster using a "virtual cluster" abstraction called Namespaces.

To begin, we'll create a Namespace called **kube-logging**. To do so, create a configuration file in your working directory (project folder) named as **kube-logging.yaml**, and write the following configuration code in that file:



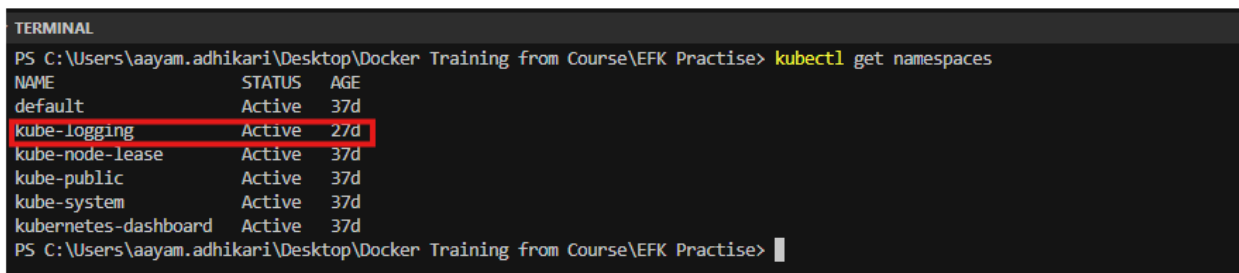
```
! kube-logging.yaml X
! kube-logging.yaml
1  apiVersion: v1
2  kind: Namespace
3  metadata:
4    name: kube-logging
```

Now, apply this configuration file using **kubectl apply -f kube-logging.yaml** in the terminal as:



```
✓ TERMINAL
PS C:\Users\aayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl apply -f kube-logging.yaml
namespace/kube-logging unchanged
PS C:\Users\aayam.adhikari\Desktop\Docker Training from Course\EFK Practise> |
```

Then, confirm if the Namespace was successfully created:



```
TERMINAL
PS C:\Users\aayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl get namespaces
NAME                STATUS   AGE
default             Active   37d
kube-logging        Active   27d
kube-node-lease     Active   37d
kube-public         Active   37d
kube-system         Active   37d
kubernetes-dashboard Active   37d
PS C:\Users\aayam.adhikari\Desktop\Docker Training from Course\EFK Practise> |
```

Step - 2: Creating the Elasticsearch Statefulset.

To start, we'll create a headless Kubernetes service called **elasticsearch** that will define a DNS domain for the 3 Pods. A headless service does not perform load balancing or have a static IP; to learn more about headless services.

To do so, create a file named **elasticsearch_svc.yaml** and write the following kubernetes service YAML:

```
! elasticsearch_svc.yaml X
! elasticsearch_svc.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: elasticsearch
5    namespace: kube-logging
6    labels:
7      app: elasticsearch
8  spec:
9    selector:
10     app: elasticsearch
11    clusterIP: None
12    ports:
13     - port: 9200
14       name: rest
15     - port: 9300
16       name: inter-node
```

Create the service using **kubectl apply -f elasticsearch_svc.yaml**, and check that the service was successfully created using **kubectl get services --namespace=kube-logging** as:

```
PS C:\Users\Aayam.Adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl get services --namespace=kube-logging
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
elasticsearch ClusterIP   None         <none>        9200/TCP,9300/TCP 27d
```

Now, create the Statefulset:

A Kubernetes StatefulSet allows you to assign a stable identity to Pods and grant them stable, persistent storage. Elasticsearch requires stable storage to persist data across Pod rescheduling and restarts.

First, open a file named **elasticsearch_stetefulset.yaml**, and write the following kubernetes yaml configuration in it:

! elasticsearch_statefulset.yaml X

! elasticsearch_statefulset.yaml

```
1  apiVersion: apps/v1
2  kind: StatefulSet
3  metadata:
4    name: es-cluster
5    namespace: kube-logging
6  spec:
7    serviceName: elasticsearch
8    replicas: 1
9    selector:
10     matchLabels:
11       app: elasticsearch
12   template:
13     metadata:
14       labels:
15         app: elasticsearch
16     spec:
17       containers:
18       - name: elasticsearch
19         image: docker.elastic.co/elasticsearch/elasticsearch:7.5.0
20         resources:
21           limits:
22             cpu: 1000m
23           requests:
24             cpu: 100m
25         ports:
26         - containerPort: 9200
27           name: rest
28           protocol: TCP
29         - containerPort: 9300
30           name: inter-node
31           protocol: TCP
```

```
containerPort: 9900
volumeMounts:
- name: data
  mountPath: /usr/share/elasticsearch/data
env:
- name: cluster.name
  value: k8s-logs
- name: node.name
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
- name: discovery.seed_hosts
  value: "es-cluster-0.elasticsearch,es-cluster-1.elasticsearch,es-cluster-2.elasticsearch"
- name: cluster.initial_master_nodes
  value: "es-cluster-0"
- name: ES_JAVA_OPTS
  value: "-Xms512m -Xmx512m"
initContainers:
- name: fix-permissions
  image: busybox
  command: ["sh", "-c", "chown -R 1000:1000 /usr/share/elasticsearch/data"]
  securityContext:
    privileged: true
  volumeMounts:
- name: data
  mountPath: /usr/share/elasticsearch/data
- name: increase-vm-max-map
  image: busybox
```

```

- name: increase-vm-max-map
  image: busybox
  command: ["sysctl", "-w", "vm.max_map_count=262144"]
  securityContext:
    privileged: true
- name: increase-fd-ulimit
  image: busybox
  command: ["sh", "-c", "ulimit -n 65536"]
  securityContext:
    privileged: true
volumeClaimTemplates:
- metadata:
  name: data
  labels:
    app: elasticsearch
  spec:
    accessModes: [ "ReadWriteOnce" ]
    # storageClassName: ""
    resources:
      requests:
        storage: 3Gi

```

Now, deploy the stateful set using command **kubectl apply -f elasticsearch_statefulset.yaml** and monitor the StatefulSet as it is rolled out using **kubectl rollout status** as follows:

```

TERMINAL powershell
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl apply -f elasticsearch_statefulset.yaml
statefulset.apps/es-cluster configured
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise>

PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl rollout status sts/es-cluster --namespace=kube-logging
partitioned roll out complete: 1 new pods have been updated...
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise>

```

Here, I've only deployed 1 pod, due to the memory limitation and also, because the minikube cluster only includes 1 node (which acts as both the master node and worker node).

Once the Pods have been deployed, you can check that your Elasticsearch cluster is functioning correctly by performing a request against the REST API.

To do so, first forward the local port `9200` to the port `9200` on one of the Elasticsearch nodes (`es-cluster-0`) using `kubectl port-forward`:

```
PS C:\Users\Aayam.Adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl port-forward es-cluster-0 9200:9200 --namespace=kube-logging
Forwarding from 127.0.0.1:9200 -> 9200
Forwarding from [::1]:9200 -> 9200
```

Now, In a separate terminal window, perform a curl request against the REST API as:

```
PS C:\Users\Aayam.Adhikari\Desktop\Docker Training from Course\EFK Practise> curl http://localhost:9200/_cluster/state?pretty

StatusCode      : 200
StatusDescription : OK
Content         : {
  "cluster_name" : "k8s-logs",
  "cluster_uuid" : "-4QdWcZYSu2QdqtEBacf5Q",
  "version"      : 216,
  "state_uuid"   : "XKYGqR4ygRCNP0yQxTj3hBw",
  "master_node"  : "U4hCwpiLRiyueMN1ClcofQ",
  "blocks"       : { }...
RawContent      : HTTP/1.1 200 OK
                  Content-Length: 376211
                  Content-Type: application/json; charset=UTF-8

                  {
                    "cluster_name" : "k8s-logs",
                    "cluster_uuid" : "-4QdWcZYSu2QdqtEBacf5Q",
                    "version"      : 216,
                    "state_uuid"   : "XKYGqR4ygRCNP0yQxTj3hBw",
                    "master_node"  : "U4hCwpiLRiyueMN1ClcofQ",
                    "blocks"       : { }...
Forms           : {}
Headers         : {[Content-Length, 376211], [Content-Type, application/json; charset=UTF-8]}
Images          : {}
InputFields     : {}
Links           : {}
ParsedHtml      : mshtml.HTMLDocumentClass
RawContentLength : 376211

PS C:\Users\Aayam.Adhikari\Desktop\Docker Training from Course\EFK Practise>
```

Step - 3: Creating the Kibana Deployment and Service.

To launch Kibana on Kubernetes, we'll create a Service called **kibana**, and a Deployment consisting of one Pod replica. You can scale the number of replicas depending on your production needs, and optionally specify a **LoadBalancer** type for the Service to load balance requests across the Deployment pods.

This time, we'll create the Service and Deployment in the same file named **kibana.yaml** as:


```
kibana.yaml X
! kibana.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: kibana
5    namespace: kube-logging
6  labels:
7    app: kibana
8  spec:
9    ports:
10   - port: 5601
11   selector:
12     app: kibana
13   ---
14   apiVersion: apps/v1
15   kind: Deployment
16   metadata:
17     name: kibana
18     namespace: kube-logging
19   labels:
20     app: kibana
21   spec:
22     replicas: 1
23     selector:
24       matchLabels:
25         app: kibana
26     template:
27       metadata:
28         labels:
29           app: kibana
```

```
template:
  metadata:
    labels:
      app: kibana
  spec:
    containers:
      - name: kibana
        image: docker.elastic.co/kibana/kibana:7.2.0
        resources:
          limits:
            cpu: 1000m
          requests:
            cpu: 100m
        env:
          - name: ELASTICSEARCH_URL
            value: http://elasticsearch:9200
        ports:
          - containerPort: 5601
```

Now, roll out the Service and Deployment using **kubectl** as follows:

```
▼ TERMINAL
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl apply -f kibana.yaml
service/kibana unchanged
deployment.apps/kibana configured
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> |
```

```
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl rollout status deployment/kibana --namespace=kube-logging
deployment "kibana" successfully rolled out
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> 
```

Now, To access the Kibana interface, we'll once again forward a local port to the Kubernetes node running Kibana. Grab the Kibana Pod details using **kubectl get** as and also forward the port **5601** to port **5601** on this pod:

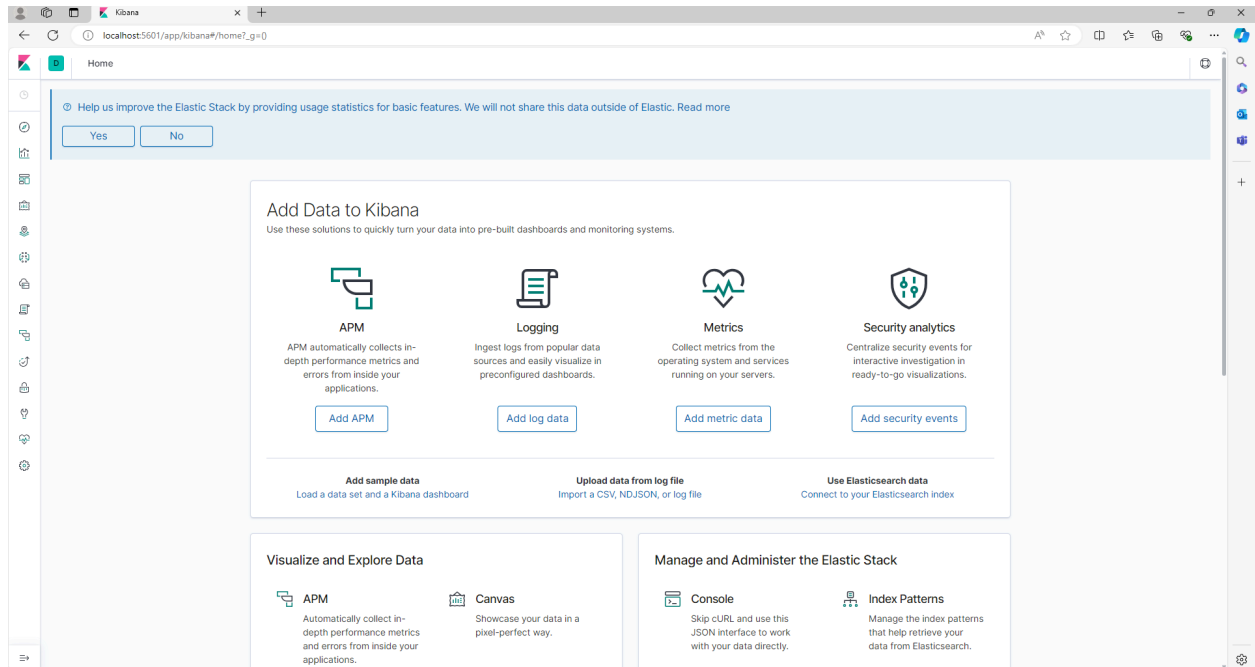
```
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl get pods --namespace=kube-logging
NAME                READY   STATUS    RESTARTS   AGE
es-cluster-0        1/1     Running   5 (16m ago)  28d
fluentd-7rz9g       1/1     Running   5 (16m ago)  28d
kibana-bc5c4875-bgkqc 1/1     Running   5 (16m ago)  28d
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl port-forward kibana-bc5c4875-bgkqc 5601:5601 --namespace=kube-logging
Forwarding from 127.0.0.1:5601 -> 5601
Forwarding from [::1]:5601 -> 5601

```

Then, In your web browser, visit the following URL:

<http://localhost:5601>

If you see the following Kibana welcome page, you've successfully deployed Kibana into your Kubernetes cluster:



Step - 4: Creating the Fluentd DaemonSet.

Create the file named **fluentd.yaml** and write the following configuration as:

! fluentd.yaml X

! fluentd.yaml

```
1  apiVersion: v1
2  kind: ServiceAccount
3  metadata:
4    name: fluentd
5    namespace: kube-logging
6    labels:
7      app: fluentd
8  ---
9  apiVersion: rbac.authorization.k8s.io/v1
10 kind: ClusterRole
11 metadata:
12   name: fluentd
13   labels:
14     app: fluentd
15 rules:
16 - apiGroups:
17   - ""
18   resources:
19     - pods
20     - namespaces
21   verbs:
22     - get
23     - list
24     - watch
25 ---
26 kind: ClusterRoleBinding
27 apiVersion: rbac.authorization.k8s.io/v1
28 metadata:
29   name: fluentd
30 roleRef:
31   kind: ClusterRole
32   name: fluentd
```

! fluentd.yaml X

! fluentd.yaml

```
29   name: fluentd
30   roleRef:
31     kind: ClusterRole
32     name: fluentd
33     apiGroup: rbac.authorization.k8s.io
34   subjects:
35   - kind: ServiceAccount
36     name: fluentd
37     namespace: kube-logging
38   ---
39   apiVersion: apps/v1
40   kind: DaemonSet
41   metadata:
42     name: fluentd
43     namespace: kube-logging
44   labels:
45     app: fluentd
46   spec:
47     selector:
48     matchLabels:
49       app: fluentd
50     template:
51     metadata:
52     labels:
53       app: fluentd
54     spec:
55       serviceAccount: fluentd
56       serviceAccountName: fluentd
57       tolerations:
58       - key: node-role.kubernetes.io/master
59         effect: NoSchedule
60     containers:
```

! fluentd.yaml ✕

! fluentd.yaml

```
46   spec:
50     template:
54       spec:
58         - key: node-role.kubernetes.io/master

60     containers:
61       - name: fluentd
62         image: fluent/fluentd-kubernetes-daemonset:v1.4.2-debian-elasticsearch
63         env:
64           - name: FLUENT_ELASTICSEARCH_HOST
65             value: "elasticsearch.kube-logging.svc.cluster.local"
66           - name: FLUENT_ELASTICSEARCH_PORT
67             value: "9200"
68           - name: FLUENT_ELASTICSEARCH_SCHEME
69             value: "http"
70           - name: FLUENTD_SYSTEMD_CONF
71             value: disable
72         resources:
73           limits:
74             memory: 512Mi
75           requests:
76             cpu: 100m
77             memory: 200Mi
78         volumeMounts:
79           - name: varlog
80             mountPath: /var/log
81           - name: varlibdockercontainers
82             mountPath: /var/lib/docker/containers
83             readOnly: true
84         terminationGracePeriodSeconds: 30
85         volumes:
86           - name: varlog
```

```
! fluentd.yaml X
! fluentd.yaml
46 spec:
50   template:
54     spec:
61       - name: fluentd
72         resources:
74           memory: 512Mi
75           requests:
76             cpu: 100m
77             memory: 200Mi
78         volumeMounts:
79           - name: varlog
80             mountPath: /var/log
81           - name: varlibdockercontainers
82             mountPath: /var/lib/docker/containers
83             readOnly: true
84         terminationGracePeriodSeconds: 30
85         volumes:
86           - name: varlog
87             hostPath:
88               path: /var/log
89           - name: varlibdockercontainers
90             hostPath:
91               path: /var/lib/docker/containers
92
```

Now, roll out the DaemonSet using kubectl:

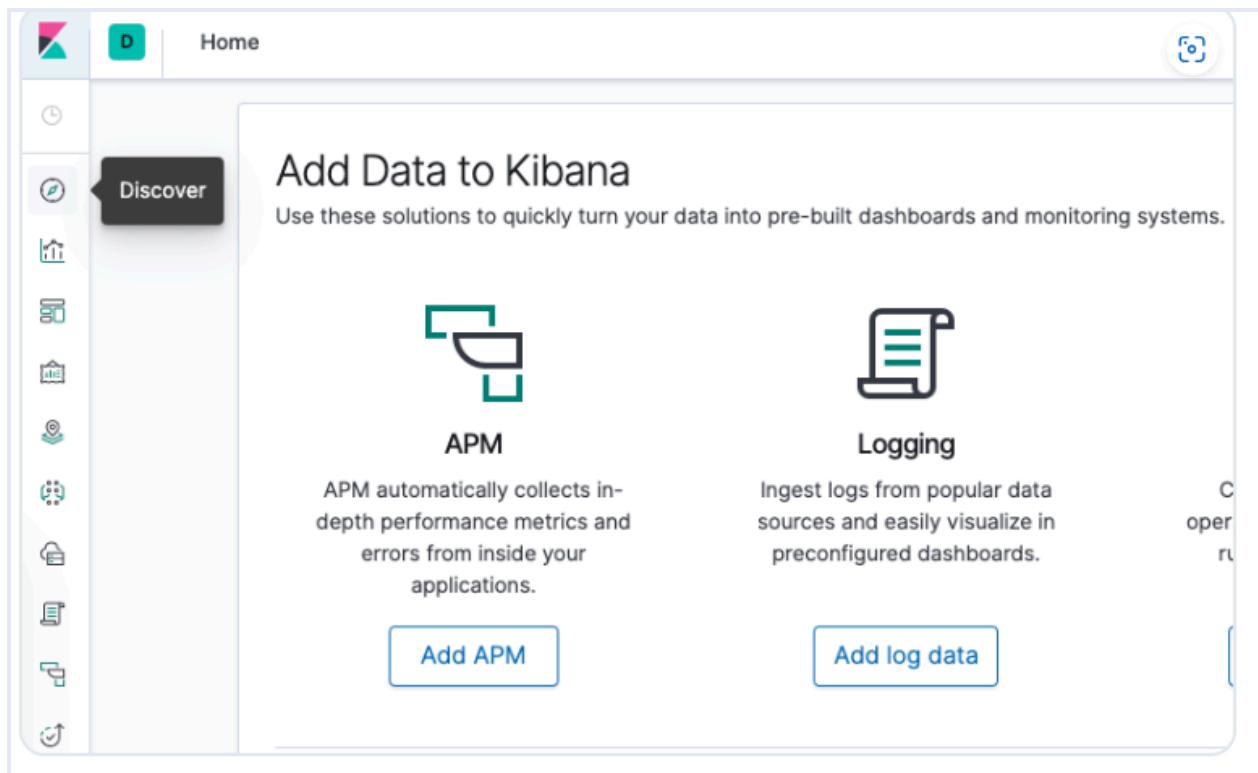
```
▼ TERMINAL
PS C:\Users\Aayam.Adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl apply -f fluentd.yaml
serviceaccount/fluentd unchanged
clusterrole.rbac.authorization.k8s.io/fluentd unchanged
clusterrolebinding.rbac.authorization.k8s.io/fluentd unchanged
daemonset.apps/fluentd unchanged
PS C:\Users\Aayam.Adhikari\Desktop\Docker Training from Course\EFK Practise> |
```

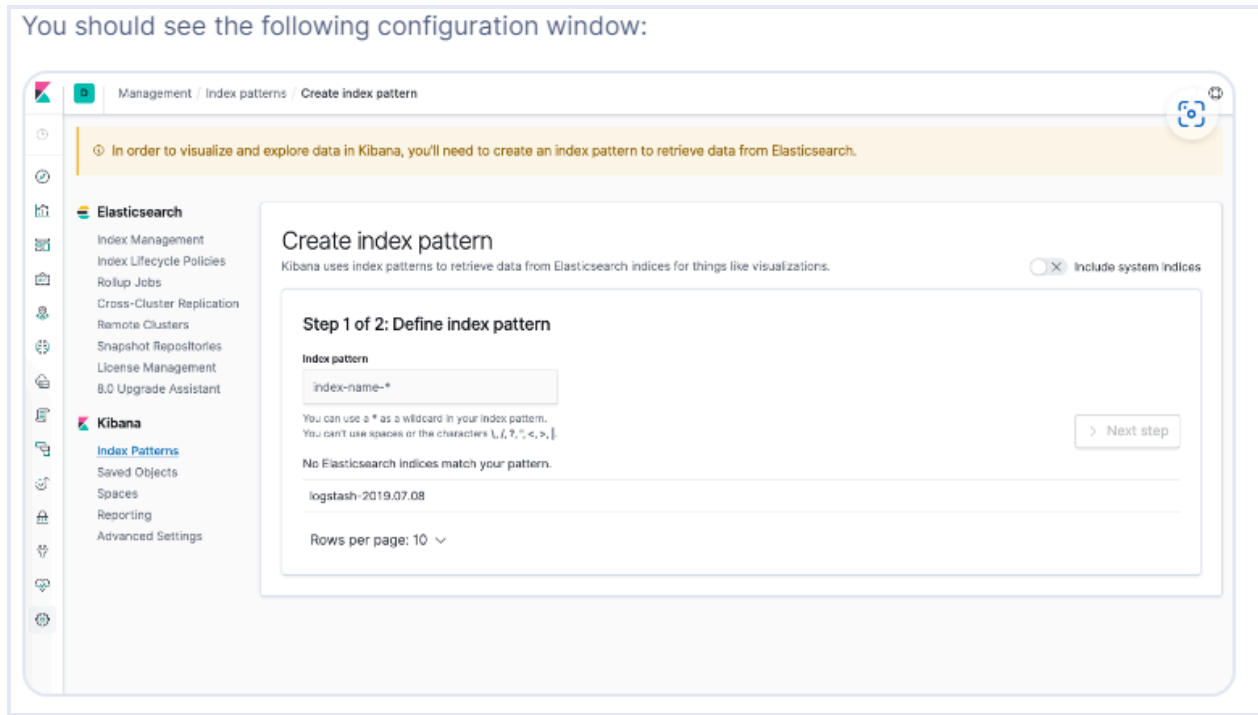
Verify that your DaemonSet rolled out successfully using **kubectl** as:

```
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl get ds --namespace=kube-logging
NAME      DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
fluentd   1         1         1       1             1           <none>          28d
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise>
```

We can now check Kibana to verify that log data is being properly collected and shipped to Elasticsearch.

With the `kubect1` port-forward still open, navigate to <http://localhost:5601> and perform the following:



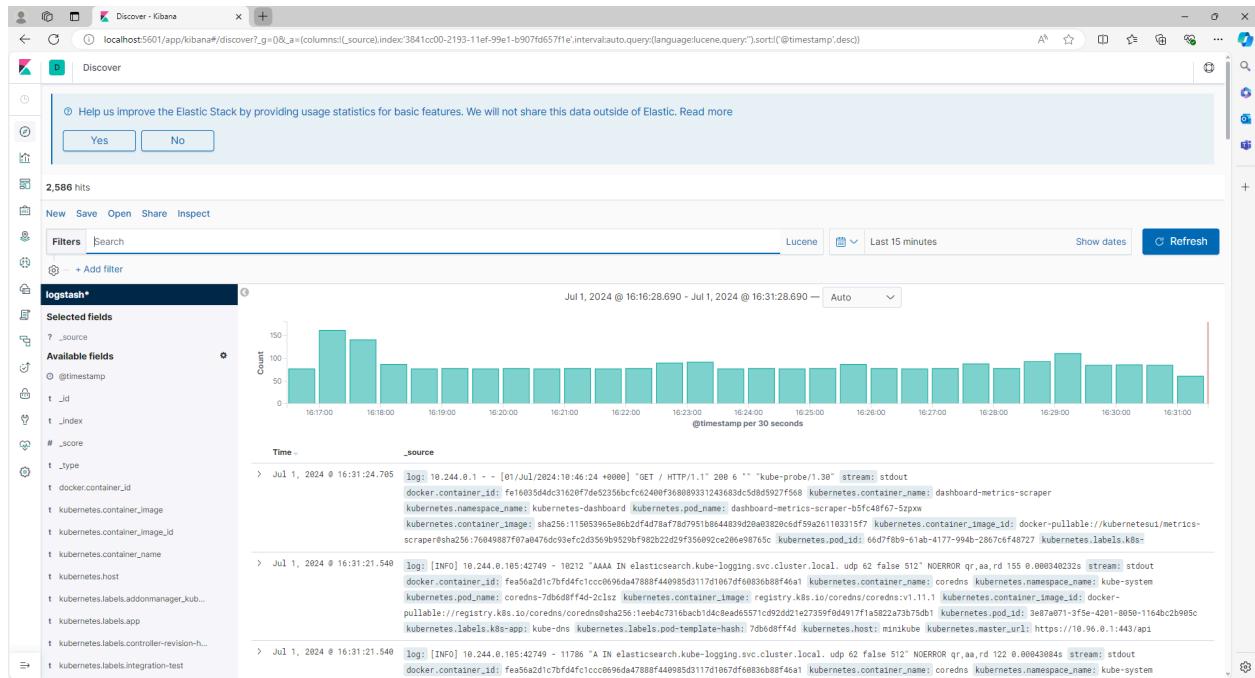


This allows you to define the Elasticsearch indices you'd like to explore in Kibana. For now, we'll just use the **logstash-*** wildcard pattern to capture all the log data in our Elasticsearch cluster. Enter **logstash-*** in the text box and click on Next step.

This allows you to configure which field Kibana will use to filter log data by time. In the dropdown, select the **@timestamp** field, and hit Create index pattern.

Now, hit Discover in the left hand navigation menu.

You should see a histogram graph and some recent log entries.



Now, for example, to see the logs of a specific **pod (or container)**, I have created a task listing project with a frontend and a backend, which is using React and Node. I have deployed this project using Kubernetes configuration files in the minikube cluster (as you can see below), and I have opened the webapp and tried to save a task as shown:

```
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl get services
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
auth-service        ClusterIP   10.109.119.80 <none>         3000/TCP         30d
frontend-service    LoadBalancer 10.99.125.1    <pending>      3000:32321/TCP   29d
kubernetes           ClusterIP   10.96.0.1      <none>         443/TCP          37d
task-service        LoadBalancer 10.108.195.151 <pending>     3000:31240/TCP   29d
users-service       LoadBalancer 10.106.97.9    <pending>     80:31616/TCP     31d

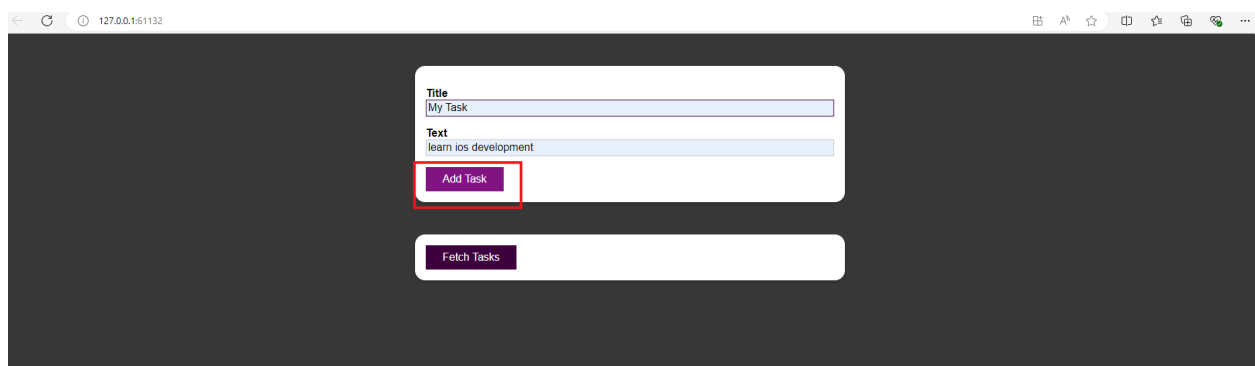
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> minikube service frontend-service
W0701 16:36:36.477986 11312 main.go:291] Unable to resolve the current Docker CLI context "default": context "default": context not found: open C:\Users\ayam.adhikari\.docker\contexts\meta\37a8eec1ce19687d132fe29051dca629d164e2c4958ba141d5f4133a33f0688f\meta.json: The system cannot find the path specified.

NAMESPACE | NAME           | TARGET PORT | URL
-----|-----|-----|-----
default | frontend-service | 3000        | http://192.168.49.2:32321

Starting tunnel for service frontend-service.

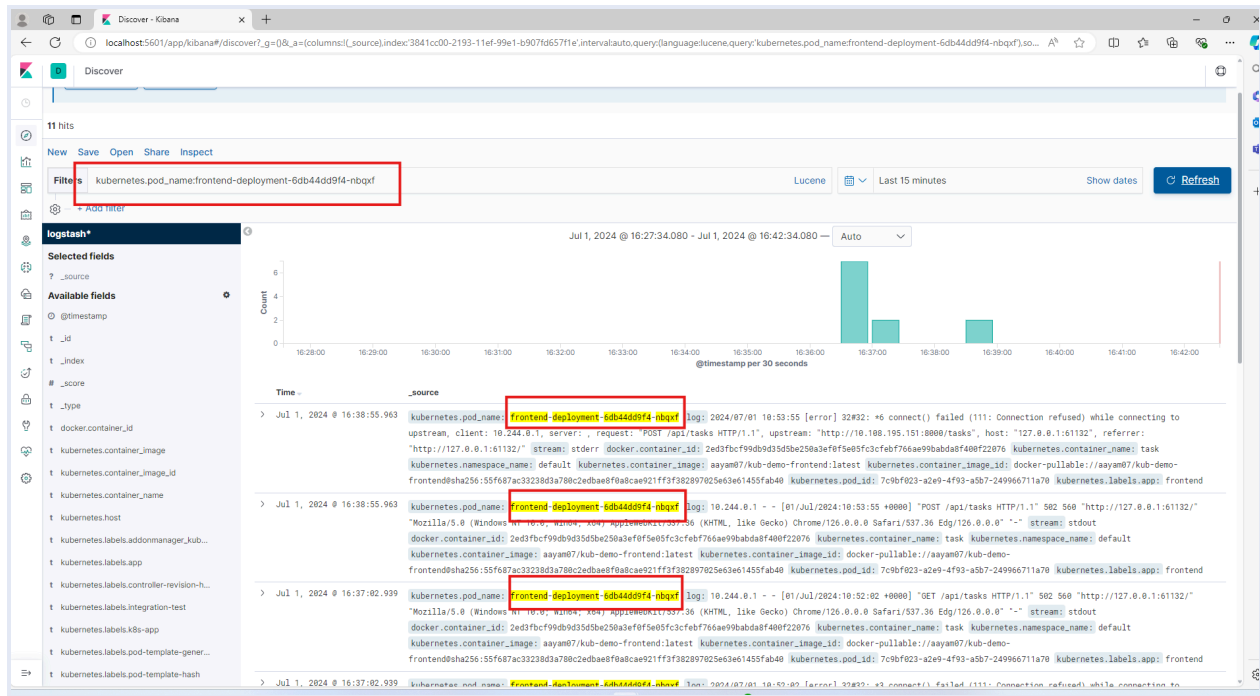
NAMESPACE | NAME           | TARGET PORT | URL
-----|-----|-----|-----
default | frontend-service |             | http://127.0.0.1:61132

Opening service default/frontend-service in default browser...
! Because you are using a Docker driver on windows, the terminal needs to be open to run it.
```



Now, navigate back to your Kibana dashboard. From the Discover page, in the search bar enter `kubernetes.pod_name:name-of-your-pod`. This filters the log data for Pods named `name-of-your-pod`. This is shown below:

```
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
auth-deployment-5d496ffd7d-2rff5    1/1     Running   3 (44m ago)  26d
frontend-deployment-6db44dd9f4-nbqxf 1/1     Running   1           171m
users-deployment-5c9b8988d6-tn517    1/1     Running   4 (44m ago)  26d
PS C:\Users\ayam.adhikari\Desktop\Docker Training from Course\EFK Practise>
```



As you can see above, we're successfully getting logs from the pod that contains the frontend React application. Similar process can be applied for any project application.